

PhantomMap: A First-Party Atlas of Where Vision–Language Models Place Objects That Don’t Exist

Author 1 (user)
Rice University
hu5@seas.upenn.edu

Author 2 (teammate)
Rice University
teammate@rice.edu

Abstract

Vision–language models such as Qwen2.5-VL-7B-Instruct [4] and LLaVA-NeXT-Mistral-7B [8] hallucinate objects absent from the input. Prior work diagnoses this either by probing internal attention [11, 5, 10] or by querying an external detector [13]. We take a complementary angle: the first-party read-out. When a VLM hallucinates, we simply ask it – through its public prompt interface – to draw a bounding box around the phantom, and analyse the resulting boxes at scale. Across 18,000 (image, object) probes on three POPE [6] splits for two models, we collect 2,806 phantom bounding boxes and show that they lie 20%–32% further from the image centre and cover 21%–37% less area than boxes drawn for real objects. A logistic regression over eight shape and logit features of the phantom box – trainable in seconds on a CPU – reaches AUROC = 0.916 on Qwen2.5-VL and 0.873 on LLaVA-NeXT, both exceeding the internal-feature HALP baseline of 0.787 [2] despite requiring nothing beyond the VLM’s API. Where a hallucination is drawn turns out to be at least as informative as what it names.

1. Introduction

Object hallucination is the most-cited failure mode of modern multimodal systems [9, 6]. Two research traditions have emerged. Internal probes (EAZY [11], OPERA [5], HalLoc [10]) open the model to analyse image-token attention or output-token confidence. External verifiers (Woodpecker [13]) forward the VLM’s caption to an object detector and flag unconfirmable entities. Both views are powerful, but neither matches what an end user actually sees.

An end user only sees the VLM’s public output. When the output asserts an object is in the image, a natural follow-up is “where, specifically, is it?” Modern grounded VLMs [4] answer with a bounding box.

When the object is absent, the bounding box is a phantom.

This paper contributes: (i) a first-party spatial atlas of 2,806 phantom and 4,971 honest boxes collected from Qwen2.5-VL-7B-Instruct and LLaVA-NeXT-Mistral-7B on the three POPE splits – phantom boxes consistently lie further from the image centre and are smaller than honest boxes, across both models; and (ii) a training-free logistic-regression detector over eight shape/location/logit features of the phantom box, reaching AUROC = 0.916 on Qwen2.5-VL and 0.873 on LLaVA-NeXT, both above the internal-feature HALP baseline of 0.787 [2]. Fig. 1 gives the flavour at a glance.




Figure 1. Top: phantom boxes that Qwen2.5-VL and LLaVA-NeXT drew around objects that are not in the image. Bottom: boxes the same two models drew for objects that are in the image. The eye already distinguishes phantoms (smaller, off-centre) from honest boxes (larger, naturally distributed).

2. Related Work

Benchmarks. POPE [6] formalised object hallucination as a yes/no existence probe on COCO [7], with

three negative-mining splits (random, popular, adversarial). AMBER [12] extends the task to multi-dimensional attribute queries. Neither asks a VLM where a hallucinated object is.

Internals and external verifiers. EAZY [11] traces hallucinated tokens back to ~ 3 high-attention image patches and zeroes them out. OPERA [5] penalises an “over-trust” attention sink at decode time. HalLoc [10] localises hallucinations at the output text-token level. All three require model internals. Woodpecker [13] chains the VLM with Grounding DINO and flags unverifiable mentions – trading one oracle for another.

Feature-based detection. HALP [2] trains a classifier on the VLM’s internal visual features and reaches AUROC = 0.787 on Qwen2.5-VL, the strongest published Qwen baseline we know of. DASH [3] clusters images that systematically trigger hallucinations. Visual CounterFact [1] studies the pixel-vs-prior conflict. Unlike all these, PhantomMap operates on the VLM’s prompted bbox output – the channel actually exposed to a user – and is the first to profile the spatial distribution of hallucinated objects.

3. Method

Elicitation prompt. For every (image I , query object o) we issue one fixed prompt (Fig. 2): “Look at the image. Is there a $\{o\}$? If yes, output $\{”bbox_2d”$: $[x1,y1,x2,y2]\}$; else output nothing.” No examples, no chain-of-thought. Outputs are parsed tolerantly (markdown fences, list-of-dict variants, LLaVA’s normalised $[0, 1]$ coordinates; §4) and clipped to $[0, W] \times [0, H]$ with minimum side 4 px. Every generated token’s greedy log-probability is logged.

Events. A phantom event is (I, o, b) where POPE labels o absent, the VLM says yes, and emits a valid bbox b . An honest event is the analogue with POPE label “yes”; honest events form the natural distribution against which phantoms are compared.

Spatial atlas. Each bbox is normalised to $[0, 1]^2$. We fit a 2D Gaussian KDE (bandwidth 0.15) over normalised centres and report the centre bias $\bar{d} = \frac{1}{N} \sum_i \sqrt{(c_{x,i} - 0.5)^2 + (c_{y,i} - 0.5)^2}$ and mean area $\bar{A} = \frac{1}{N} \sum_i w_i h_i$.

Training-free detector. For every “yes” event with a valid bbox we extract eight features: the normalised centre (c_x, c_y) , width w , height h , area wh , centre distance d_{center} , aspect ratio, and the mean and minimum

log-probability of the four coordinate tokens (plus the yes/no token’s log-probability). Labels are 1 for phantoms, 0 for honest. After standardising, we fit an ℓ_2 -regularised logistic regression ($C=1$) and report AUROC on a stratified 70/30 split. Baselines: chance, the single-feature logit-only score $-\log p_{\text{yes}}$, and the published HALP internal-feature AUROC of 0.787 [2].

4. Experiments and Results

Setup. Qwen2.5-VL-7B-Instruct [4] and LLaVA-NeXT-Mistral-7B [8], both in bf16 on a single NVIDIA L4 GPU via Hugging Face transformers with greedy decoding. Inputs: POPE random / popular / adversarial – each 3,000 (image, object) probes on COCO val2014 [7], 9,000 probes per model. A full sweep for both models completes in ~ 6 h. Qwen emits pixel coordinates; LLaVA emits normalised $[0, 1]$ and grounds on only 19% of its positive responses. The parser handles both formats with no model-specific code.

Atlas (Fig. 3, Table 2). Both models show the same pattern: phantom boxes have a larger mean distance-to-centre (0.314 vs 0.261 for Qwen; 0.304 vs 0.231 for LLaVA) and a smaller mean area (0.106 vs 0.167 for Qwen; 0.189 vs 0.239 for LLaVA). Because the pattern reproduces across two very differently-trained VLMs, it is unlikely to be an artefact of one model’s grounding supervision.

Detector (Fig. 4, Table 3). The full-feature LR reaches AUROC = 0.916 on Qwen ($n=6,965$) and 0.873 on LLaVA ($n=812$), both exceeding HALP’s 0.787. On Qwen the bbox geometry adds to the yes-token confidence ($0.899 \rightarrow 0.916$); on LLaVA, geometry and confidence overlap and are near-indistinguishable.

Analysis. Qwen grounds on virtually every positive (100%) while LLaVA grounds on only 19%, so LLaVA’s yes-token logprob already captures most of the separating signal and geometry adds little. Qualitatively, phantoms cluster into three patterns (Fig. 5): oversized hedge boxes that half-cover the scene; peripheral boxes pressed to a corner; and edge-clipped boxes flush with an image border.

5. Conclusion

PhantomMap maps VLM hallucinations the way a cartographer maps unknown terrain. Using only the VLM’s public bbox interface on POPE, we show that phantoms are placed with a predictable spatial signature – smaller, further from centre – across two 7B-scale VLMs with very different grounding recipes, and




Figure 2. PhantomMap pipeline. For each (I, o) pair in POPE, the VLM is prompted for yes/no plus a bbox. When POPE labels the object absent, every “yes” is a phantom event and its bbox feeds both the atlas (§3) and the detector (§3).

Model	Split	n	Acc.	Hall. rate	Yes rate	Valid bbox%
Qwen2.5-VL-7B	adv.	3000	0.658	0.625	0.783	100.0
Qwen2.5-VL-7B	pop.	3000	0.683	0.575	0.758	100.0
Qwen2.5-VL-7B	rnd.	3000	0.659	0.623	0.782	100.0
LLaVA-NeXT-7B	adv.	3000	0.860	0.145	0.506	19.1
LLaVA-NeXT-7B	pop.	3000	0.898	0.070	0.468	18.9
LLaVA-NeXT-7B	rnd.	3000	0.918	0.029	0.448	19.1

Table 1. POPE accuracy, hallucination rate, yes rate, and the fraction of “yes” answers that yielded a parseable bbox. Qwen is more willing to ground and more willing to hallucinate; LLaVA is the opposite.




Figure 3. Phantom atlas. KDE of phantom-box centres vs honest-box centres, per model. Phantom boxes are displaced outward relative to the honest distribution for both models. Colour = density; white dots = individual events.

that this signature feeds a training-free detector that exceeds the strongest internal-feature baseline. Where a hallucination is drawn turns out to carry at least as much information as what it names, and it is accessible to anyone with API access to the model. Because the same features could help adversaries hide hallucinations, we release code under an academic-use licence

and add no new images to the public domain; POPE inherits COCO’s biases and so does our atlas.

Acknowledgments. We used Claude Opus 4.7 (Anthropic) for early brainstorming, literature-scan sanity checks, and boilerplate-code drafting; every experimental choice, run, and figure is our own. We thank the authors of POPE [6], Qwen2.5-VL [4], LLaVA-NeXT

Model	Set	n	$\bar{d}_{\text{center}}$	$\bar{A}$
Qwen2.5-VL-7B	phantom	2735	0.314	0.106
Qwen2.5-VL-7B	honest	4230	0.261	0.167
LLaVA-NeXT-7B	phantom	71	0.304	0.189
LLaVA-NeXT-7B	honest	741	0.231	0.239

Table 2. Atlas summary. Phantoms are further from the image centre and smaller in area than honest boxes, for both models.




Figure 4. Detector ROC. Training-free logistic regression on phantom-box features (solid) vs. a logit-only baseline (dashed) vs. chance. Axes are false-positive and true-positive rates.

Detector	Qwen2.5-VL-7B	LLaVA-NeXT-7B
Random	0.500	0.500
HALP (published)[2]	0.787	—
Logit only	0.899	0.889
LR full (ours)	0.916	0.873

Table 3. Hallucination-detection AUROC on stratified 70/30 splits over “yes” events that yielded a valid bbox ($n=6,965$ Qwen; $n=812$ LLaVA). Our first-party detector exceeds HALP’s internal-feature detector on Qwen by +12.8 AUROC points.




Figure 5. Failure-mode taxonomy. Top: phantoms covering most of the image. Middle: phantoms pressed against the periphery. Bottom: phantoms flush with an image border.

[8], and EAZY [11]. Author 1 wrote the evaluation code and ran all VLM experiments; Author 2 built the atlas visualisations, produced the figures, and wrote the related work.

References

- [1] Anonymous. Pixels versus priors: Controlling knowledge priors in vision-language models through visual counterfactuals. arXiv preprint arXiv:2505.17127, 2025.
- [2] Anonymous. HALP: Detecting hallucinations in vision-language models without generating a single token. arXiv preprint arXiv:2603.05465, 2026.
- [3] M. Augustin et al. DASH: Detection and assessment of systematic hallucinations of VLMs. In IEEE/CVF International Conference on Computer Vision (ICCV), 2025.
- [4] S. Bai, K. Chen, X. Liu, J. Wang, W. Ge, S. Song, K. Dang, P. Wang, S. Wang, J. Tang, et al. Qwen2.5-VL Technical Report. arXiv preprint arXiv:2502.13923, 2025.
- [5] Q. Huang, X. Dong, P. Zhang, B. Wang, C. He, J. Wang, D. Lin, W. Zhang, and N. Yu. OPERA: Alleviating hallucination in multi-modal large language models via over-trust penalty and retrospection-allocation. In IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR), 2024.
- [6] Y. Li, Y. Du, K. Zhou, J. Wang, W. X. Zhao, and J.-R. Wen. Evaluating object hallucination in large vision-language models. In Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing (EMNLP), 2023.
- [7] T.-Y. Lin, M. Maire, S. Belongie, J. Hays, P. Perona, D. Ramanan, P. Dollár, and C. L. Zitnick. Microsoft COCO: Common objects in context. In European Conference on Computer Vision (ECCV), 2014.
- [8] H. Liu, C. Li, Y. Li, B. Li, Y. Zhang, S. Shen, and Y. J. Lee. LLaVA-NeXT: Improved reasoning, OCR, and world knowledge. In Blog post, llava-vl.github.io, 2024.

- [9] H. Liu, W. Xue, Y. Chen, D. Chen, X. Zhao, K. Wang, L. Hou, R. Li, and W. Peng. A survey on hallucination in large vision-language models. arXiv preprint arXiv:2402.00253, 2024.
- [10] J. Park et al. HalLoc: Token-level localization of hallucinations for vision language models. In IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR), 2025.
- [11] Pseudoc18 et al. EAZY: A training-free approach on detecting and mitigating object hallucinations in LVLMs through hallucinatory image tokens. In Technical report, pseudoc18.github.io/eazyVLM, 2025. read: project page and method section.
- [12] J. Wang, Y. Wang, G. Xu, J. Zhang, Y. Gu, H. Jia, M. Yan, J. Zhang, and J. Sang. AMBER: An LLM-free multi-dimensional benchmark for MLLMs hallucination evaluation. arXiv preprint arXiv:2311.07397, 2023.
- [13] S. Yin, C. Fu, S. Zhao, T. Xu, H. Wang, D. Sui, Y. Shen, K. Li, X. Sun, and E. Chen. Woodpecker: Hallucination correction for multimodal large language models. Science China Information Sciences, 2024.